

EDAN96

Applied Machine Learning

Lecture 8: Classification Techniques and Neural Networks

Pierre Nugues

`Pierre.Nugues@cs.lth.se`

November 23, 2022

Content

Overview and practice of some neural network architectures:

- Datasets
- Input formatting
- Feedforward networks

We will use:

- PyTorch, <https://pytorch.org>, a powerful API to design and train network, and
- scikit-learn, <https://scikit-learn.org/stable/>, a general purpose machine-learning toolkit.


Goal of the Lecture

Describe a language detector: Given a string predict the language:

- *Bonjour* → French
- Guten Tag → German

Follow the complete compact language detector (CLD3)

<https://github.com/google/cld3>



You will implement it during the fourth lab

Dataset Exploration

When dealing with a dataset, the first step is to explore it.

All machine-learning experiments should start with it.

Such a step is called *exploratory data analysis* (EDA)

It consists of measuring:

- the size of the dataset (observations)
- the number of classes, number of observations per class
- the mean and standard deviation of the variables
- the domain of categorical classes
- number of missing values, etc.

EDA and Visualization

Many EDAs show results on graphics

Many examples on <https://www.kaggle.com/>, for instance

[https://www.kaggle.com/code/imoore/](https://www.kaggle.com/code/imoore/intro-to-exploratory-data-analysis-eda-in-python)

[intro-to-exploratory-data-analysis-eda-in-python](https://www.kaggle.com/code/imoore/intro-to-exploratory-data-analysis-eda-in-python) and

[https://www.kaggle.com/code/ekami66/](https://www.kaggle.com/code/ekami66/detailed-exploratory-data-analysis-with-python)

[detailed-exploratory-data-analysis-with-python](https://www.kaggle.com/code/ekami66/detailed-exploratory-data-analysis-with-python)

The Tatoeba Dataset

- A language detector (lab 4) is a kind of categorization
- Categorization is by far the most popular application of machine learning
- See here for the datasets: <https://huggingface.co/tasks>,
<https://archive.ics.uci.edu/ml/datasets.php>
- As dataset to train CLD3, we will use Tatoeba
<https://tatoeba.org/>

Code Example

Experiment: An elementary EDA Jupyter Notebook:

https://github.com/pnugues/edan96/blob/main/programs/5-tatoeba_eda_select.ipynb

Dataset Formatting

When training a model, a common practice is to split the dataset in three sets:

- 1 Train, validation (also called development), and test sets
- 2 Already done in many datasets, see <https://huggingface.co/datasets>
- 3 Otherwise do it. You will have to do it for CLD3

This partitioning is essential to measure overfit when training. Identical test sets give means to compare results and performance of systems (see scientific papers).

Code Example

Experiment: An elementary EDA Jupyter Notebook:

https://github.com/pnugues/edan96/blob/main/programs/5-tatoeba_eda_select.ipynb

Input Formatting

Starting from observations, we format the input to use it as input to a classifier

This step is also called the vectorization: Turn the input into numerical vectors

Digitization: We will assume all our data is in digital format

Numerical input: We already have vectors or tensors to represent our data. We have to standardize them

Categorical or text input: Linear classifiers do not accept this: We have to encode the symbols as vectors, using one-hot encoding (also called indicator vector or function)

Hashing: One-hot encoding may result in very large vectors, millions of parameters, when dealing with text. We can reduce their size with hashing techniques

Embeddings: One-hot encoding is a sparse representation. We can design network architectures to produce smaller (dim=100) and dense vectors

Standardization

Most algorithms in classification do not work well if the parameters have widely differing ranges:

- 1 Parameter 1: [10,000,000; 20,000,000]
- 2 Parameter 2: [-0.01; 0.002]

Before training a model, we must standardize the data.

Two options:

- 1 Remove the mean and divide by the standard deviation with respect to columns:

$$x_{i,jstd} = \frac{x_{i,j} - \bar{x}_{\cdot,j}}{\sigma_{x_{\cdot,j}}}.$$

- 2 Norm the rows to 1:

$$x_{i,jnorm} = \frac{x_{i,j}}{\sqrt{\sum_{k=0}^{n-1} x_{i,k}^2}}.$$

Fisher's Iris Dataset (1936)

180 MULTIPLE MEASUREMENTS IN TAXONOMIC PROBLEMS

Table I

<i>Iris setosa</i>				<i>Iris versicolor</i>				<i>Iris virginica</i>			
Sepal length	Sepal width	Petal length	Petal width	Sepal length	Sepal width	Petal length	Petal width	Sepal length	Sepal width	Petal length	Petal width
5.1	3.5	1.4	0.2	7.0	3.2	4.7	1.4	6.3	3.3	6.0	2.5
4.9	3.0	1.4	0.2	6.4	3.2	4.5	1.5	5.8	2.7	5.1	1.9
4.7	3.2	1.3	0.2	6.9	3.1	4.9	1.5	7.1	3.0	5.9	2.1
4.6	3.1	1.5	0.2	5.5	2.3	4.0	1.3	6.3	2.9	5.6	1.8
5.0	3.6	1.4	0.2	6.5	2.8	4.6	1.5	6.5	3.0	5.8	2.2
5.4	3.9	1.7	0.4	5.7	2.8	4.5	1.3	7.6	3.0	6.6	2.1
4.6	3.4	1.4	0.3	6.3	3.3	4.7	1.6	4.9	2.5	4.5	1.7
5.0	3.4	1.5	0.2	4.9	2.4	3.3	1.0	7.3	2.9	6.3	1.8
4.4	2.9	1.4	0.2	6.6	2.9	4.6	1.3	6.7	2.5	5.8	1.8
4.9	3.1	1.5	0.1	5.2	2.7	3.9	1.4	7.2	3.6	6.1	2.5
5.4	3.7	1.5	0.2	5.0	2.0	3.5	1.0	6.5	3.2	5.1	2.0
4.8	3.4	1.6	0.2	5.9	3.0	4.2	1.5	6.4	2.7	5.3	1.9
4.8	3.0	1.4	0.1	6.0	2.2	4.0	1.0	6.8	3.0	5.5	2.1
4.3	3.0	1.1	0.1	6.1	2.9	4.7	1.4	5.7	2.5	5.0	2.0
5.8	4.0	1.2	0.2	5.6	2.9	3.6	1.3	5.8	2.8	5.1	2.4
5.7	4.4	1.5	0.4	6.7	3.1	4.4	1.4	6.4	3.2	5.3	2.3
5.4	3.9	1.3	0.4	5.6	3.0	4.5	1.5	6.5	3.0	5.5	1.8
5.1	3.5	1.4	0.3	5.8	2.7	4.1	1.0	7.7	3.8	6.7	2.2
5.7	3.8	1.7	0.3	6.2	2.2	4.5	1.5	7.7	2.6	6.9	2.3
5.1	3.8	1.5	0.3	5.6	2.5	3.9	1.1	6.0	2.2	5.0	1.5
5.4	3.4	1.7	0.2	5.9	3.2	4.8	1.8	6.9	3.2	5.7	2.3
5.1	3.7	1.5	0.4	6.1	2.8	4.0	1.3	5.6	2.8	4.9	2.0

Code Example

Experiment: sklearn and its cancer dataset with a Jupyter Notebook:
<https://github.com/pnugues/edan96/blob/main/programs/6-standardization.ipynb>

Categorical Values

Linear classifiers only understand numbers

A collection of two documents D1 and D2:

D1: Chrysler plans new investments in Latin America.

D2: Chrysler plans major investments in Mexico.

How to represent these words or these documents?

Categorical Values: One-hot encoding

Let us suppose we want to predict the parts of speech of the words, for instance *plans* in:

Input:	Chrysler	<i>plans</i>	major	investments	in	Mexico
Output:	PNoun	Verb	Adjective	Noun	Prep.	PNoun

One-hot encoding:

- 1 Assigns a unique index to each symbol. The number of indices corresponds to the number of symbols:
 $\{\text{'Chrysler': 1, 'in': 2, 'investments': 3, 'major': 4, 'Mexico': 5, 'plans': 6}\}$
- 2 Represent a symbol of index i by a vector $\mathbf{x} = (x_1, x_2, \dots, x_n)$, where n is the largest index and all the coordinates are 0, except $x_i = 1$

'Chrysler': (1, 0, 0, 0, 0, 0)

'in': (0, 1, 0, 0, 0, 0)

'investments': (0, 0, 1, 0, 0, 0)

...

Categorical Values: Multi-hot encoding

A collection of two documents D1 and D2:

D1: Chrysler plans new investments in Latin America.

D2: Chrysler plans major investments in Mexico.

Multi-hot encoding (also called a bag-of-words representation):

- 1 Creates a index of all the symbols (words) in all the documents
- 2 For each document, creates a set of its symbols (word)
- 3 Represents a document by a vector of 0s and 1s. $x_i = 1$ if the word of index i is in the document, or 0 otherwise.

D.	america	chrysler	in	investments	latin	major	mexico	new	plans
1	1	1	1	1	1	0	0	1	1
2	0	1	1	1	0	1	1	0	1

Hashing

One-hot encoding leads to very large dimensions as there are billions of different words.

In the Tatoeba experiment (see notebook), the number of unigrams, bigrams, and trigrams is (4844, 88471, 294472)

A solution is to hash the symbols and use the remainder of a division (modulo) as index

```
>>> hash('abc')  
-6712881850779232724  
>>> hash('abc') % 100  
76    # Always less than 100
```

Hashing:

- 1 Reduces the vectors size and makes it manageable
- 2 Creates conflicts: two symbols can have the same hash numbers
- 3 Is usable in classification

Code Example

Experiment: hashing CLD3 n-grams Jupyter Notebook:

https://github.com/pnugues/edan96/blob/main/programs/7-ngram_hashing.ipynb

Reproducible Hash Codes

```
def reproducible_hash(string):  
    """  
    reproducible hash on any string  
  
    Arguments:  
        string: python string object  
  
    Returns:  
        signed int64  
    """  
  
    # We are using MD5 for speed not security.  
    h = hashlib.md5(string.encode("utf-8"),  
                     usedforsecurity=False)  
    return int.from_bytes(h.digest()[0:8], 'big', signed=True)
```

Dense Vectors

We can replace one-hot vectors by dense ones using embeddings
Many techniques, often based on language modeling, here CBOW
Guess a missing word given its context. Using the example:

Sing, O **goddess**, the anger of Achilles son of Peleus,

Cloze test: A reader, given the incomplete phrase:

Sing, O ____, the anger of Achilles

has to fill in the blank with the correct word, here **goddess**.

Easy to create a dataset for a Cloze test

$$\mathbf{X} = \begin{bmatrix} \text{sing} & \text{o} & \text{the} & \text{anger} \\ \text{o} & \text{goddess} & \text{anger} & \text{of} \\ \text{goddess} & \text{the} & \text{of} & \text{achilles} \\ \text{the} & \text{anger} & \text{achilles} & \text{son} \\ \text{anger} & \text{of} & \text{son} & \text{of} \\ \text{of} & \text{achilles} & \text{of} & \text{peleus} \end{bmatrix}; \mathbf{y} = \begin{bmatrix} \text{goddess} \\ \text{the} \\ \text{anger} \\ \text{of} \\ \text{achilles} \\ \text{son} \end{bmatrix}$$


CBOW Architecture

Dimensionality reduction inside a neural network.

Context words one-hot encoded followed by an **embedding layer**.

A dense representation is a trainable vector of 10 to 300 dimensions.

The vector parameters are learned in the fitting procedure.



Embeddings in PyTorch

PyTorch has an `Embedding(num_embeddings, embedding_dim, ...)` class

An embedding object is a matrix from which we can extract embedding vectors using an index

This is just a lookup table

```
# Creates trainable vectors of size 64
embedding_chars = nn.Embedding(MAX_CHARS, 64)
```

```
# Extracts embeddings in rows 3 and 2,
# corresponding to two characters
embedding_chars(torch.LongTensor([3, 2]))
```


Code Example

Experiment: Embeddings with a Jupyter Notebook:

https://github.com/pnugues/edan96/blob/main/programs/8-pytorch_embeddings.ipynb

Sum of Embeddings in CLD3

CLD3 computes the weighted sum of embeddings



Embedding Bags in PyTorch

EmbeddingBags class creates embedding objects.

Given a list of embeddings (a list of rows) as input, an embedding bag returns the weighted sum of the embeddings.

We specify the weights with a `per_sample_weights` parameter.

```
embedding_bag = nn.EmbeddingBag(MAX_CHARS, 64, mode='sum')
```

```
# Computes the sum of rows 1 and 2 and rows 3 and 4
```

```
# The result is a matrix of two rows
```

```
embedding_bag(torch.tensor([[1, 2], [3, 4]]))
```

```
embedding_bag(torch.tensor([[1, 2], [3, 4]]),  
               per_sample_weights=torch.tensor([[0.5, 0.5],  
                                                [0.2, 0.8]]))
```

Embedding Bags in PyTorch (II)

With bags of inequal sizes, we have to use a list of offsets


```
embedding_bag(torch.tensor([1, 2, 3, 4]),  
              offsets=torch.tensor([0, 2]),  
              per_sample_weights=torch.tensor([0.5, 0.5, 0.2, 0.8]))
```

Code Example

Experiment: Embedding bags with a Jupyter Notebook:


https://github.com/pnugues/edan96/blob/main/programs/8-pytorch_embeddings.ipynb

Neural Networks




A photomicrograph showing the classic view of the snail-shaped cochlea with hair cells stained green and neurons showing reddish-purple. [Decibel Therapeutics (<https://www.decibeltx.com>)]. Source: <https://www.genengnews.com/insights/targeting-the-inner-ear/>

Logistic Regression as a Neural Network




Neural Networks




Activation Functions

The perceptron




Logistic regression




Activation Functions

Rectified linear unit (ReLU), where

$$\text{ReLU}(x) = \max(0, x).$$



Neural Networks with Hidden Layers



A Text Dataset: *Salammbô*

A corpus is a collection – a body – of texts.

French original


English translation

GUSTAVE FLAUBERT

SALAMMBÔ

ÉDITION DÉFINITIVE
AVEC DES DOCUMENTS NOUVEAUX

PARIS
G. CHARPENTIER, ÉDITEUR
12, RUE DE CHEVREUIL-SAINT-GERMAIN, 12
—
1883
Tous droits réservés



Classification Dataset

Dataset for binary classification: *Salammô* in French (1) and English (0)

	# char.	# A	class (y)	# char.	# A	class (y)
Chapter 1	36,961	2,503	1	35,680	2,217	0
Chapter 2	43,621	2,992	1	42,514	2,761	0
Chapter 3	15,694	1,042	1	15,162	990	0
Chapter 4	36,231	2,487	1	35,298	2,274	0
Chapter 5	29,945	2,014	1	29,800	1,865	0
Chapter 6	40,588	2,805	1	40,255	2,606	0
Chapter 7	75,255	5,062	1	74,532	4,805	0
Chapter 8	37,709	2,643	1	37,464	2,396	0
Chapter 9	30,899	2,126	1	31,030	1,993	0
Chapter 10	25,486	1,784	1	24,843	1,627	0
Chapter 11	37,497	2,641	1	36,172	2,375	0
Chapter 12	40,398	2,766	1	39,552	2,560	0
Chapter 13	74,105	5,047	1	72,545	4,597	0
Chapter 14	76,725	5,312	1	75,352	4,871	0
Chapter 15	18,317	1,215	1	18,031	1,119	0

Code Example

Experiment: Jupyter Notebook:

https://github.com/pnugues/edan96/blob/main/programs/10-Salamambo_multi_torch-2022.ipynb